

Class08

Kaliyah Adei-Manu

Background

In today's class we will apply the emthods and techniques clustering and PCA to help make sense of real world breast cancer FNA biopsy data set

Data import

We will start by importikng our data. It is a CSV file so we will use the `read.csv()` function.

```
fna.data <- "https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names=1)
```

We will have a look at the first few rows

```
head(wisc.df)
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean
842302	M	17.99	10.38	122.80	1001.0
842517	M	20.57	17.77	132.90	1326.0
84300903	M	19.69	21.25	130.00	1203.0
84348301	M	11.42	20.38	77.58	386.1
84358402	M	20.29	14.34	135.10	1297.0
843786	M	12.45	15.70	82.57	477.1
	smoothness_mean	compactness_mean	concavity_mean	concave.points_mean	
842302	0.11840	0.27760	0.3001	0.14710	
842517	0.08474	0.07864	0.0869	0.07017	
84300903	0.10960	0.15990	0.1974	0.12790	
84348301	0.14250	0.28390	0.2414	0.10520	
84358402	0.10030	0.13280	0.1980	0.10430	

843786	0.12780	0.17000	0.1578	0.08089		
	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302	0.2419		0.07871	1.0950	0.9053	8.589
842517	0.1812		0.05667	0.5435	0.7339	3.398
84300903	0.2069		0.05999	0.7456	0.7869	4.585
84348301	0.2597		0.09744	0.4956	1.1560	3.445
84358402	0.1809		0.05883	0.7572	0.7813	5.438
843786	0.2087		0.07613	0.3345	0.8902	2.217
	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	153.40	0.006399		0.04904	0.05373	0.01587
842517	74.08	0.005225		0.01308	0.01860	0.01340
84300903	94.03	0.006150		0.04006	0.03832	0.02058
84348301	27.23	0.009110		0.07458	0.05661	0.01867
84358402	94.44	0.011490		0.02461	0.05688	0.01885
843786	27.19	0.007510		0.03345	0.03672	0.01137
	symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302	0.03003		0.006193	25.38	17.33	
842517	0.01389		0.003532	24.99	23.41	
84300903	0.02250		0.004571	23.57	25.53	
84348301	0.05963		0.009208	14.91	26.50	
84358402	0.01756		0.005115	22.54	16.67	
843786	0.02165		0.005082	15.47	23.75	
	perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302	184.60	2019.0		0.1622	0.6656	
842517	158.80	1956.0		0.1238	0.1866	
84300903	152.50	1709.0		0.1444	0.4245	
84348301	98.87	567.7		0.2098	0.8663	
84358402	152.20	1575.0		0.1374	0.2050	
843786	103.40	741.6		0.1791	0.5249	
	concavity_worst	concave.points_worst	symmetry_worst			
842302	0.7119		0.2654	0.4601		
842517	0.2416		0.1860	0.2750		
84300903	0.4504		0.2430	0.3613		
84348301	0.6869		0.2575	0.6638		
84358402	0.4000		0.1625	0.2364		
843786	0.5355		0.1741	0.3985		
	fractal_dimension_worst					
842302		0.11890				
842517		0.08902				
84300903		0.08758				
84348301		0.17300				
84358402		0.07678				
843786		0.12440				

Make sure to remove the first `diagnosis` column - I don't want to use this for my machine learning models. We will use it later to compare our results to the expert diagnosis.

```
wisc.data <- wisc.df[,-1]
diagnosis <- wisc.df$diagnosis
```

Q1. How many observations are in this dataset?

```
nrow(wisc.data)
```

[1] 569

Q2. How many of the observations have a malignant diagnosis?

```
table(diagnosis)
```

```
diagnosis
  B   M
357 212
```

212 have a malignant diagnosis

Q3. How many variables/features in the data are suffixed with `_mean`?

```
colnames <- colnames(wisc.data)
length(grep("_mean", colnames))
```

[1] 10

There are 10 columns that are suffixed `_mean`

Principal Component Analysis

The main function here is `prcomp()` and we want to make sure we set the optional argument `scale=TRUE`

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731
Proportion of Variance	0.00314	0.00266	0.00198	0.00175	0.00165	0.00104	0.0010
Cumulative Proportion	0.98649	0.98915	0.99113	0.99288	0.99453	0.99557	0.9966
	PC22	PC23	PC24	PC25	PC26	PC27	PC28
Standard deviation	0.16565	0.15602	0.1344	0.12442	0.09043	0.08307	0.03987
Proportion of Variance	0.00091	0.00081	0.0006	0.00052	0.00027	0.00023	0.00005
Cumulative Proportion	0.99749	0.99830	0.9989	0.99942	0.99969	0.99992	0.99997
	PC29	PC30					
Standard deviation	0.02736	0.01153					
Proportion of Variance	0.00002	0.00000					
Cumulative Proportion	1.00000	1.00000					

Q4. From your results, what proportion of the original variance is captured by the first principal component (PC1)?

The first principal component captures 44.27% of the variance.

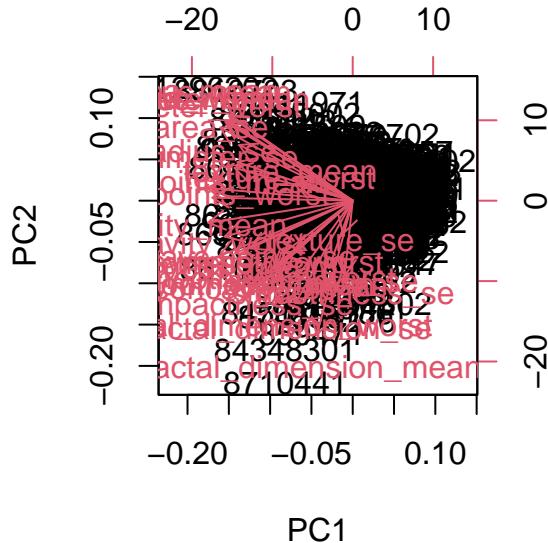
Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

3 PC's are required to describe 70% of the original variance, the 3rd PC cumulatively captures 72.6% of the original variance.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

7 PC's are needed to capture at least 90% of the original variance. PCA 7 collectively accounts for 91.0% of the original variance.

```
biplot(wisc.pr)
```



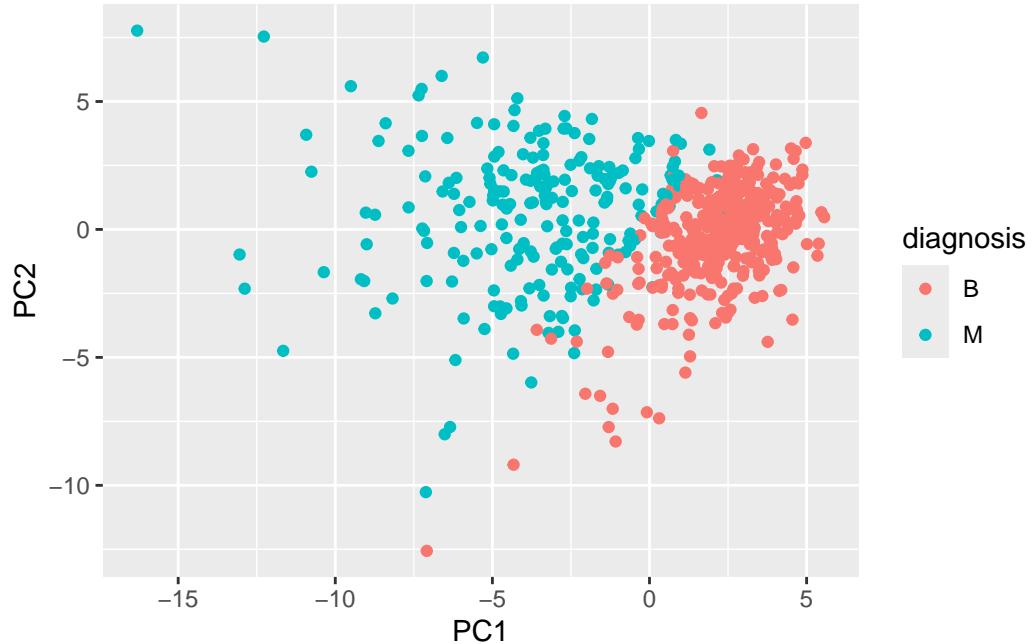
Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

Its quite difficult to read a lot of the information and data overlaps making it difficult to interpret.

Our main PCA “score plot” or “PC plot” of results:

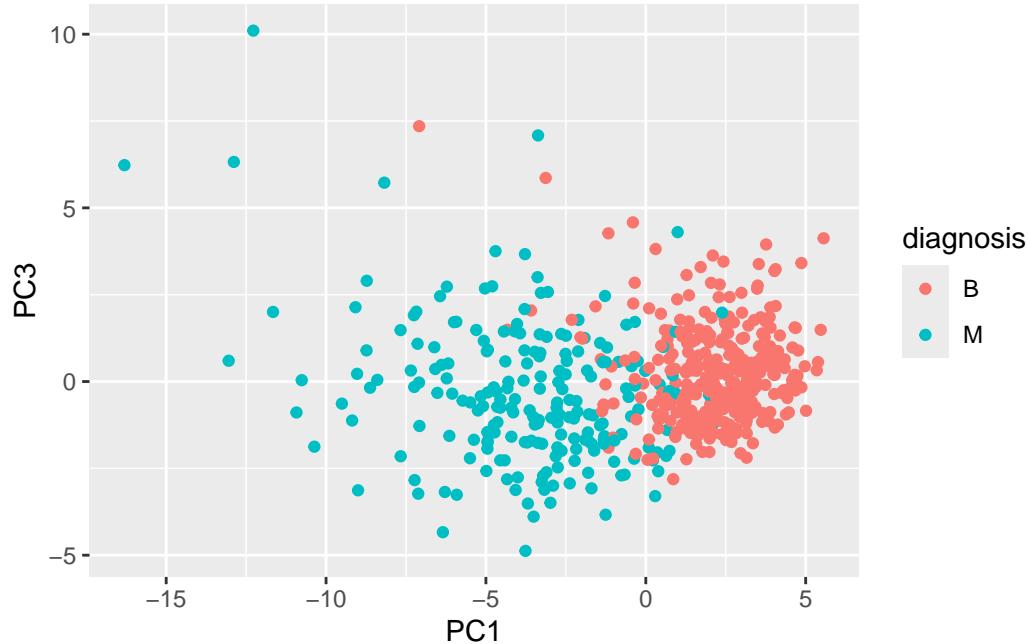
```
library(ggplot2)
```

```
ggplot(wisc.pr$x)+  
  aes(PC1, PC2, col=diagnosis) +  
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x) +
  aes(PC1, PC3, col=diagnosis) +
  geom_point()
```



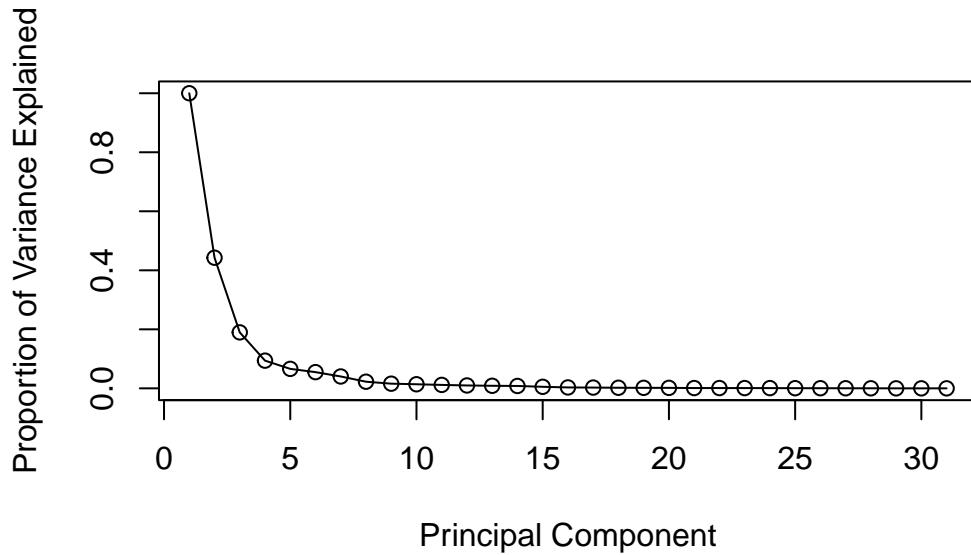
This plot has more overlap of the two clusters than the plot of PC1 vs PC2, this plot doesn't appear to have a linear relationship. ## Interpreting PCA results

Variance Explained

```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

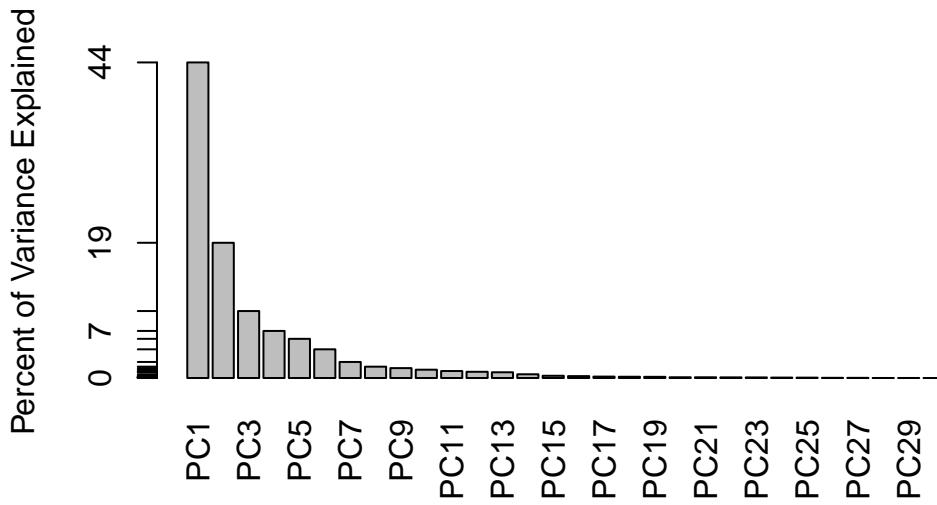
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var/30
plot(c(1,pve), xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



An Alternative Scree Plot

```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```



Communicating PCA Results:

Q9. For the first principal component, what is the component of the loading vector (i.e. `wisc.pr$rotation[,1]`) for the feature `concave.points_mean`? This tells us how much this original feature contributes to the first PC. Are there any features with larger contributions than this one?

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean
-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean	radius_se	texture_se
-0.06436335	-0.20597878	-0.01742803
perimeter_se	area_se	smoothness_se
-0.21132592	-0.20286964	-0.01453145
compactness_se	concavity_se	concave.points_se
-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se	radius_worst
-0.04249842	-0.10256832	-0.22799663

<code>texture_worst</code>	<code>perimeter_worst</code>	<code>area_worst</code>
-0.10446933	-0.23663968	-0.22487053
<code>smoothness_worst</code>	<code>compactness_worst</code>	<code>concavity_worst</code>
-0.12795256	-0.21009588	-0.22876753
<code>concave.points_worst</code>	<code>symmetry_worst</code>	<code>fractal_dimension_worst</code>
-0.25088597	-0.12290456	-0.13178394

The ‘concave.point_mean value is -0.26085376, this shows how much this contributes to PCA 1

Hierachial Clustering

First scale the data (with the `scale()` function), then calculate a distance matrix (with the `dist()` function). Then cluster with the `hclust()` function and plot:

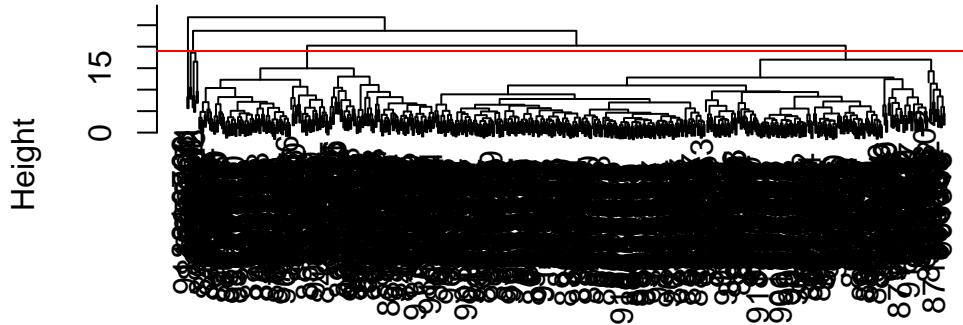
```
wisc.hclust <- hclust(dist(scale(wisc.data)))
```

Q10. Using the `plot()` and `abline()` functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty= 2)
```

Warning in int_abline(a = a, b = b, h = h, v = v, untf = untf, ...): "lyt" is not a graphical parameter

Cluster Dendrogram



```
dist(scale(wisc.data))
hclust (*, "complete")
```

You can also use the `cutree()`

```
wisc.hclust.clusters <- cutree(wisc.hclust, k=4)
table(wisc.hclust.clusters)
```

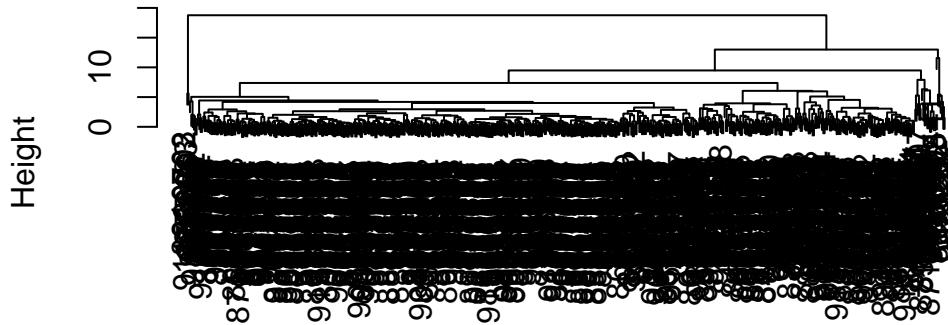
```
wisc.hclust.clusters
 1   2   3   4
177   7 383   2
```

Combining methods

here we will take our PCA results and use those as input for clustering. In other words our `wisc.pr$x` scores that we plotted above (the main output from PCA - how the data lies on our new principal component axis/variables) and use a subset of these PCs that capture the most variance as input for `hclust()`

```
pc.dist <- dist(wisc.pr$x[,1:3])
wisc.pr.hclust <- hclust(pc.dist, method="average")
plot(wisc.pr.hclust)
```

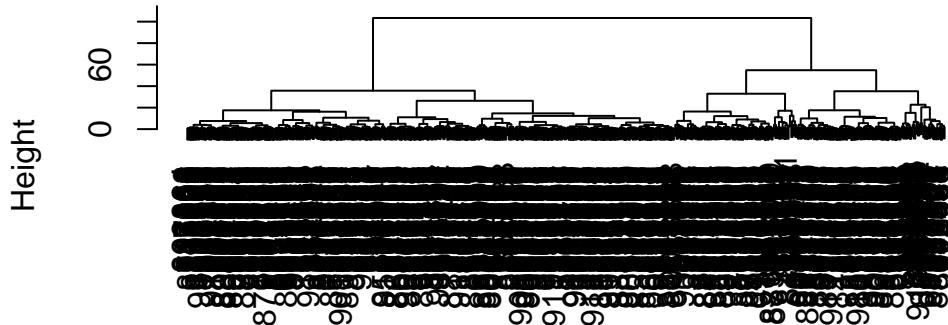
Cluster Dendrogram



```
pc.dist  
hclust (*, "average")
```

```
pc.dist <- dist(wisc.pr$x[,1:3])  
wisc.pr.hclust <- hclust(pc.dist, method="ward.D2")  
plot(wisc.pr.hclust)
```

Cluster Dendrogram



```
pc.dist  
hclust (*, "ward.D2")
```

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

I liked “ward.D2” , for this dendrogram I think it made more clear clusters then “average”

Let’s find out if these two clusters are malignant and benign

Cut the dendrogram/ tree into two main groups/ clusters:

```
grps <- cutree(wisc.pr.hclust, k=2)
table(grps)
```

```
grps
  1   2
203 366
```

I want to know how many of the clustering grps with values of 1 or 2 correspond to the expert diagnosis

Q. Q13. How well does the newly created hclust model with two clusters separate out the two “M” and “B” diagnoses?

```
table(grps,diagnosis)
```

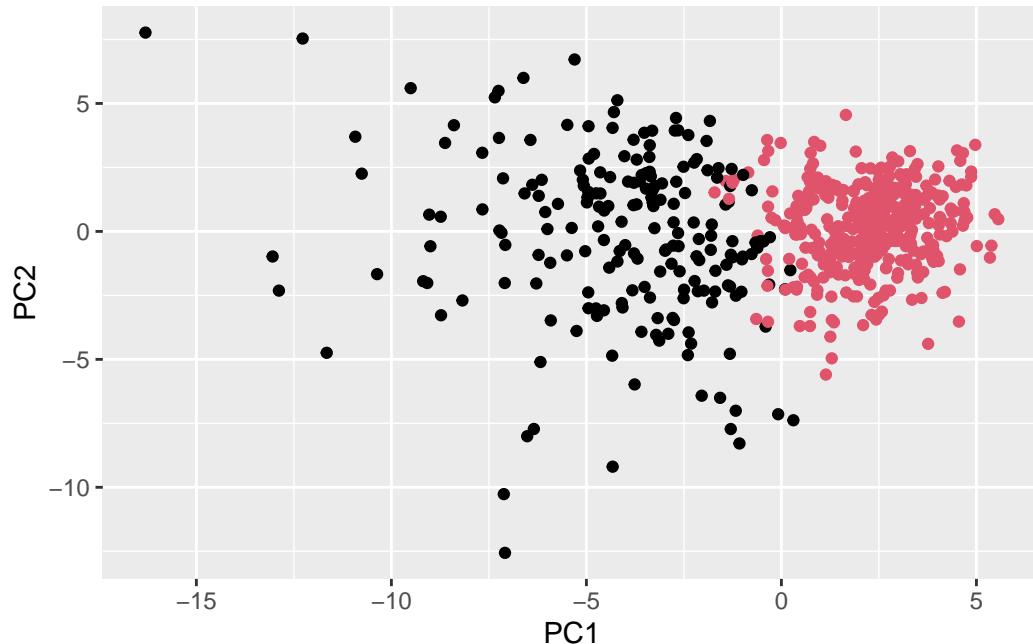
```
diagnosis
grps   B   M
 1   24 179
 2 333  33
```

Q14. How well do the hierarchical clustering models you created in the previous sections (i.e. without first doing PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.hclust.clusters and wisc.pr.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

```
diagnosis
wisc.hclust.clusters   B   M
 1   12 165
 2   2   5
 3 343  40
 4   0   2
```

```
ggplot(wisc.pr$x) +
  aes(PC1, PC2) +
  geom_point(col=grps)
```



My clustering **group 1** are mostly “M” diagnosis (179) and my clustering **group 2** are mostly “B” diagnosis

24 FP 179 TP 333 TN 33 FN

Sensitivity $TP/(TP+FN)$

179/(179+33)

[1] 0.8443396

Specificity $TN/(TN+FP)$

333/(333+24)

[1] 0.9327731

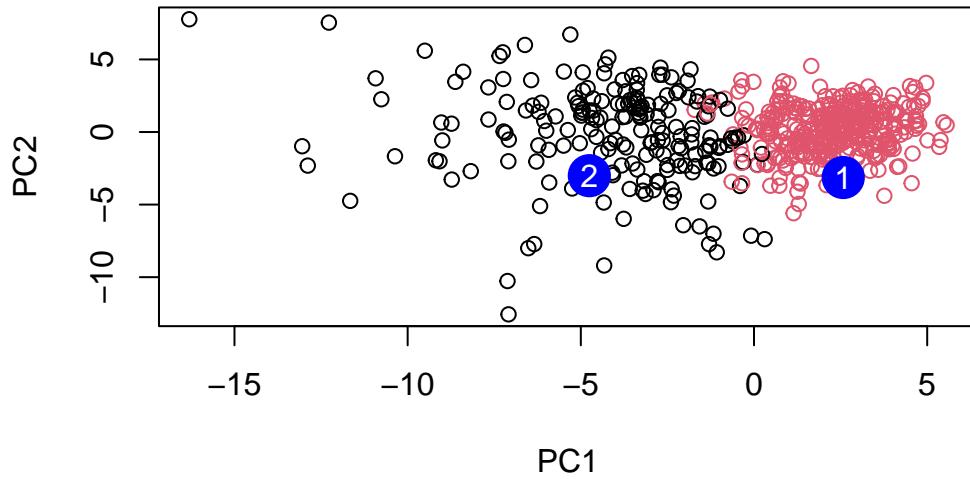
Use this PCA model to predict using data from Michigan:

```
url <- "https://tinyurl.com/new-samples-CSV"
new <- read.csv(url)
npc <- predict(wisc.pr, newdata=new)
npc
```

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
[1,]	2.576616	-3.135913	1.3990492	-0.7631950	2.781648	-0.8150185	-0.3959098
[2,]	-4.754928	-3.009033	-0.1660946	-0.6052952	-1.140698	-1.2189945	0.8193031
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
[1,]	-0.2307350	0.1029569	-0.9272861	0.3411457	0.375921	0.1610764	1.187882
[2,]	-0.3307423	0.5281896	-0.4855301	0.7173233	-1.185917	0.5893856	0.303029
	PC15	PC16	PC17	PC18	PC19	PC20	
[1,]	0.3216974	-0.1743616	-0.07875393	-0.11207028	-0.08802955	-0.2495216	
[2,]	0.1299153	0.1448061	-0.40509706	0.06565549	0.25591230	-0.4289500	
	PC21	PC22	PC23	PC24	PC25	PC26	
[1,]	0.1228233	0.09358453	0.08347651	0.1223396	0.02124121	0.078884581	
[2,]	-0.1224776	0.01732146	0.06316631	-0.2338618	-0.20755948	-0.009833238	
	PC27	PC28	PC29	PC30			
[1,]	0.220199544	-0.02946023	-0.015620933	0.005269029			
[2,]	-0.001134152	0.09638361	0.002795349	-0.019015820			

Making a new PCA plot for Michigan :

```
plot(wisc.pr$x[,1:2], col=grps)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")
```



Q16. Which of these new patients should we prioritize for follow up based on your results?

We should prioritize analyzing the patients that overlap. In other words prioritize follow up for patients where there are both black and red dots, meaning these patients are more likely to be misdiagnosed than the patients who are closer to the centers of their clusters.